Homework 2: Shell Implementation

Due Date: November 1, 2021, at 11:55 PM

Academic Honesty

Aside from the narrow exception for collaboration on homework, all work submitted in this course must

be your own. Cheating and plagiarism will not be tolerated. If you have any questions about a specific

case, please ask the instructor or the TAs. Also, we will be testing for plagiarism. NYU's Policy on

Academic Misconduct:

https://engineering.nyu.edu/campus-and-community/student-life/office-student-affairs/policies/student-co
de-conduct#chapter-id-34265

Homework Notes

General Notes

● Read the assignment carefully, including what files to include.

● Don't assume limitations unless they are explicitly stated.

● Treat provided examples as just that, not an exhaustive list of cases that should work.

● When in doubt regarding what needs to be done, ask. Another option is to test it in the real UNIX

operating system. Does it behave the same way?

● **TEST your solutions, make sure they work. It's obvious when you didn't test the code.**

● Remember if you are working with a partner, submit your name and NetID# along with your

   partner's name and NetID# along with your other documents on Brightspace.

Notes

● This assignment does not involve modifying or using xv6 (although the code for Shell.c is adapted from the xv6 shell).

● You must write, compile, and test your code on the Unix shell in Anubis (although this code won't be running in xv6 this time).

● Because we're not running shell.c in xv6, you are free to use the standard C libraries such as stdlib.h, string.h, stdio.h, etc.

● While it may be tempting to just copy xv6's implementation, there are enough differences between the xv6 APIs and those in Anubis's Unix OS that doing so would be a bad idea. You can look at how it works for inspiration though.

In this assignment, you will implement pieces of the UNIX shell and get some familiarity with a few UNIX library calls along with the UNIX process model. By the end of the assignment, you will have implemented a shell that can run any series of complex pipelines of commands, such as:

```
$ cat words.txt | grep cat | sed s/cat/dog/ > doggerel.txt
```

The pipeline shown above takes a word document labelled words.txt (a file generally installed on UNIX systems that contains a list of English words. This file will be included to download on Brightspace), select the words containing the string "cat", and then uses sed to replace "cat" with "dog", so that, for example, "concatenate" becomes "condogenate". The results are then outputted to a text file labelled "*doggerel.txt*". (You can find detailed descriptions of each of the commands in the pipeline by consulting the manual page for the command; e.g.: "*man grep*" or "*man sed*".)

Start by downloading the shell.c and words.txt. Shell.c is a skeleton file attached to this homework which you will place in your Anubis IDE. You don't have to understand how the parser works in detail, but you should have a general idea of how the flow of control works in the program. You will also see comments labelled with "*//your code here"*, which is where you will implement the functionality to make the shell actually work. Next, you must try to compile the source code to the Anubis Unix shell (Again, this will not be using xv6):

```
$ gcc shell.c -o shell
```

You can then run and interact with the shell by typing `./shell` :

```
user@cs6233:~$ ./shell

cs6233> ls

exec not implemented

cs6233>
```

Note: The command prompt for our shell is set to `cs6233>` to make it easy to tell the difference between our shell and the Anubis's Linux shell. You can quit your shell by typing *Ctrl-C* or *Ctrl-D*

Problem 1: Command Execution (15 points)

Implement basic command execution by filling in the code inside of the `case''` blocking the `runcmd` function. You will want to look at the manual page for the exec(3) function by typing `"man 3 exec"` (Note: throughout this course, when referring to commands that one can look up in the man pages, we will typically specify the section number in parentheses -- thus, since exec is found in section 3, we will say exec(3)).

Once this is done, you should be able to use your shell to run single commands, such as

```
cs6233> ls

cs6233> grep cat

words.txt
```

Hint:

You will notice that there are many variants on exec(3). You should read through the differences between them in the xv6 manual, and then choose the one that allows you to run the commands above -- in particular, pay attention to whether the version of exec you're using requires you to enter in the full path to the program, or whether it will search the directories in the PATH environment variable.

Make sure to explain your program or approach in a few sentences and submit your explanation in a text document on Brightspace (5 points)

Problem 2: I/O Redirection (15 points)

Now extend the shell to handle input and output redirection. Programs will be expecting their input on standard input and write to standard output, so you will have to open the file and then replace standard input or output with that file. As before, the parser already recognizes the '>' and '<' characters and builds a `redircmd` structure for you, so you just need to use the information in that `redircmd` to open a file and replace standard input or output with it.

Hints:

1. Look at the dup2(2) and open(2) calls.

2. The file descriptor the program is currently using for input or output is available in `rcmd->fd`.

3. If you're confused about where `rcmd->fd` is coming from, look at the `redircmd` function and remember that 0 is standard input, 1 is standard output.

4. Be careful with the `open` call; in particular, make sure you read about the case when you pass the `O_CREAT` flag.

When this is done, you should be able to redirect the input and output of commands:

```
cs6233> ls > a.txt

cs6233> sort -r < a.txt
```

Make sure to explain your program or approach in a few sentences and submit your explanation in a text document on Brightspace (5 points)

Problem 3: Pipes (20 points)

The final task is to add the ability to pipe the output of one command into the input of another.

You will fill out the code for the '|' case of the switch statement in `runcmd` to do this.

Hints:

1. The parser provides the left command in `pcmd->left` and the right command in `pcmd- >right`.

2. Look at the fork(2), pipe(2), close(2) and wait(2) calls.

3. If your program just hangs, it may help to know that reads to pipes with no data will block until all file descriptors referencing the pipe are closed.

4. Note that fork(2) creates an exact copy of the current process. The two processes share any file descriptors that were open at the time the fork occurred. You can get a sense for this:

```
#include <stdio.h>

#include <unistd.h>

#include <fcntl.h>

#include <sys/stat.h>

int main() {

int filedes;

filedes = open("myfile.txt", O_RDWR | O_CREAT, S_IRUSR | S_IWUSR);

int rv;

rv = fork();

if (rv == 0) {

    char msg[] = "Process 1\n";

    printf("Hello, I'm in the child, my process ID is %d\n",getpid());

    write(filedes, msg, sizeof(msg));

}

else {
```

```
    char msg[] = "Process 2\n";

    printf("This is the parent process, my process ID is %d and my

child is %d\n", getpid(), rv);

    write(filedes, msg, sizeof(msg));

}


    close(filedes);

}
```

If you put that code into a separate file in Anubis, compile it, and then run the resulting program, you should see a result like:

```
anubis@anubis-ide:~/homework2_fall2021-cs6233student$:~$ gcc a.c -o
a.out

anubis@anubis-ide:~/homework2_fall2021-cs6233student$:~$ ./a.out

This is the parent process, my process ID is [Random Number X] and my
child is [Random Number X + 1]
Hello, I'm in the child, my process ID is [Random Number X + 1]

anubis@anubis-ide:~/homework2_fall2021-cs6233student$:~$ cat
myfile.txt

Process 2
Process 1
```

You can see that both the parent and child process both got a copy of "filedes", and that writes to it from each process went to the same underlying file.

5. You may find it helpful to re-read the first chapter of the xv6 manual, which describes in detail how

the xv6 shell works. Note that the code shown there will not work as-is -- you will have to adapt it for

the Anubis IDE Unix environment.

Once this is done, you should be able to run a full pipeline:

```
cs6233> cat words.txt | grep cat | sed s/cat/dog/ >

doggerel.txt

cs6233> grep con < doggerel.txt
```

Explain your program or approach in a few sentences (10 points)

Explain in detail: -

Q1. How would you implement lists of commands, separated by ";" - (10 points)

Q2. How would you implement sub shells by implementing "(" and ")" - (10 points)

Q3. How would you implement running commands in the background by supporting "&" and "wait" – (10 points)

- You may submit your modified shell.c in Anubis. In addition, you must also submit a pdf/word file to the Brightspace assignment containing your written answers.

Rubric

Please remember that these are only **some** of the cases we'll be testing for. Make sure to test your program thoroughly and thoughtfully.

Total: 100 points

- Program does not compile (-90)

- Stdin redirection does not work (-30)

- Stdout redirection does not work (-30)

- Pipe does not work (-40)

- Exec does not work (-30)

- Bad permissions on redirected output (-10)

- "*Not implemented*" messages left in (-5)

- Missing your explanation on implementing shell command executions (-5)

- Missing your written explanation on implemented shell output redirection (-5)

- Missing your written explanation on implementing shell pipe redirection (-5)

- Missing your written answers regarding how you would implement shell's unique parameters (-10 for each of the three questions)

Credits: This assignment is adapted from a homework assignment by Brendan Dolan-Gavitt